

Author: Zara Irfan

GitHub: [Zara-Irfan](#)

## Gradient Descent for Linear Regression

### Overview

This Python script demonstrates the implementation of **Gradient Descent** to perform simple linear regression on a small dataset. The goal is to fit a line that best approximates the relationship between the independent variable  $x$  and the dependent variable  $y$  by iteratively minimizing the mean squared error.

The code visualizes the learning process by plotting the regression line at each iteration, showing how the line converges towards the best fit.

### Features

- Implements gradient descent from scratch without using external ML libraries.
- Visualizes each step of the gradient descent process.
- Final plot shows the original data points with the optimized regression line.
- Adjustable hyperparameters: learning rate ( $\alpha$ ) and number of iterations.

### Code Explanation

- **Data:** The dataset consists of five points with input values  $x$  and corresponding target values  $y$ .
- **Initialization:** The slope ( $m$ ) and intercept ( $b$ ) are initialized to zero.
- **Gradient Descent Loop:** For a fixed number of iterations, the script:
  - Predicts the outputs ( $y_{\text{pred}}$ ) based on current parameters.
  - Calculates the gradients of the loss function with respect to  $m$  and  $b$ .
  - Updates  $m$  and  $b$  by moving against the gradient scaled by the learning rate.
  - Plots the current regression line (with transparency) to visualize convergence.
- **Final Visualization:** Displays the original data points in red and the converged regression line.

### Usage

1. Clone or download this repository.
2. Ensure you have the required libraries installed: numpy and matplotlib.
3. Run the script in a Jupyter notebook or Python environment.
4. Observe the stepwise improvement of the regression line toward fitting the data.

bash

CopyEdit

```
pip install numpy matplotlib
```

python

CopyEdit

```
# Run the provided Python code to see the gradient descent in action
```

```
python gradient_descent_linear_regression.py
```

### **Code**

python

CopyEdit

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
# Sample data
```

```
x = np.array([1, 2, 3, 4, 5])
```

```
y = np.array([1.5, 3.7, 4.1, 6.0, 8.2])
```

```
# Hyperparameters
```

```
alpha = 0.01 # learning rate
```

```
iterations = 50
```

```
m = 0 # initial slope
```

```
b = 0 # initial intercept
```

```
plt.figure()

# Gradient Descent Loop
for _ in range(iterations):
    y_pred = m * x + b
    error = y - y_pred

    # Plot current line (in green)
    plt.plot(x, y_pred, color='green', alpha=0.3)

    # Gradient calculation
    m_grad = -(2 / len(x)) * sum(x * error)
    b_grad = -(2 / len(x)) * sum(error)

    # Update parameters
    m -= alpha * m_grad
    b -= alpha * b_grad

# Final plot with data points
plt.scatter(x, y, color='red') # actual data
plt.title("Gradient Descent for Linear Regression")
plt.xlabel("x")
plt.ylabel("y")
plt.show()
```

Gradient Descent for Linear Regression

